Deep Learning in Audio, Homework 1: Automatic Speech Recognition

Alexey Slizkov
26.10.2023

I have implemented the steps of the instruction, also beam search and beam search with LM (using 3rd party lib). The metrics of BSwLM are significantly better than those without.

I have a total of 176 wanbd runs, though most of them were interrupted early. Here is the summary of the most successful ones (BS&LM metrics not present here, only plain model):

#	num	hidden	params	lr	bs	steps	time	train			valid			test-clean			test-other			comments
#	n_ls	dim					/ep.	loss	cer	wer	loss	cer	wer	loss	cer	wer	loss	cer	wer	
Baseline model:																				
94		4096	17.4M	3e-4	64	15x1000	03:40	3.68	.898	1.00	3.96	.910	1.00	3.91	.900	1.00	4.12	.909	1.00	I thought I was running RNN...
95		4096	17.4M	3e-3	64	10x1000	03:40	2.41	.943	.999	2.37	.936	.999	2.35	.938	.999	2.59	.960	.999	-- same --
RNN:																				
99	3	4096	84.5M	3e-3	64	25x100	01:36	31.3	.988	.998	21.5	.986	.989	21.4	.985	.985	22.1	.984	.984	exploded
switched to linear_lr, time per epoch is now between starts of two trains																				
101	3	4096	84.5M	3e-4	64	47x100	03:00	1.19	.387	.845	1.12	.374	.839	1.10	.369	.832	1.58	.495	.912	
104	3	4096	84.5M	1e-3	64	50x100	03:00	1.08	.355	.808	.979	.327	.788	.959	.320	.776	1.42	.451	.882	
LSTM:																				
108	3	2048	85.0M	1e-3	32	50x100	03:01	.789	.255	.658	.680	.224	.617	.675	.221	.607	1.15	.361	.786	not sure that 03:01
117	6	2048	185M	1e-3	24	7x100	09:15	2.88	.988	.996	2.91	.985	.982	2.90	.985	.983	2.92	.983	.982	cer&wer were 1 up to 5th epoch, .99999 on 6th
121	6	1024	46.7M	1e-3	48	38x100	01:26	2.79	.990	.999	2.82	.985	.999	2.82	.984	.999	2.84	.983	.999	same for 1-5 epochs
126	3	4096	337M	1e-3	48	11x100	14:09	2.25	.671	1.08	2.26	.691	1.05	2.26	.685	.700	2.36	.700	1.04	duration matches #108
127	3	2048	85.0M	1e-3	32	363x100	04:14	.735	.233	.607	.645	.208	.575	.632	.202	.562	1.15	.351	.760	scheduler stopped at 364:37 instead of 367 epochs
Bidirectional LSTM:																				
132	3	1024	59.8M	1e-3	48	50x100	01:28	.425	.134	.397	.347	.110	.352	.349	.110	.349	.790	.241	.593	record
137	3	1024	63.9M	1e-3	256	50x100		3.20	.959	.987	3.22	.933	.983	3.23	.931	.984	3.18	.921	.984	added 3 conv1d layers before LSTMs, kernel size 5, stride 2
142	3	1024	63.9M	1e-3	48	50x100	01:33	.464	.146	.411	.379	.120	.361	.378	.118	.353	.841	.255	.593	here stride 1, padding 2
143	3	1024	61.1M	1e-3	48	50x100	01:32	.418	.130	.376	.332	.104	.322	.335	.104	.321	.775	.236	.562	only 2 convolutions
146	3	1024	68.1M	1e-3	48	50x100	01:39	.395	.124	.361	.328	.103	.320	.329	.102	.317	.773	.233	.558	3 convs, all intermediates fed into LSTMs
Implemented BeamSearch, not using it																				
154	3	1024	68.1M	1e-3	48	15x100	01:36	.777	.246	.629	.680	.216	.598	.672	.213	.586	1.13	.349	.764	
155	3	1024	67.7M	1e-3	48	50x100	01:36	.418	.130	.379	.347	.110	.338	.349	.108	.335	.783	.236	.567	kernel size 5->3, hence padding 1
167	3	1024	89.2M	1e-3	32	50x100	01:51	2.83	.999	1.00	2.81	1.00	1.00	2.90	1.00	1.00	2.91	1.00	1.00	convs are now 2d 3x3
170	3	1024	67.7M	1e-3	48	30x100	01:36	.567	.178	.489	.468	.147	.439	.465	.146	.431	.917	.279	.645	1d convs are back
171	3	1024	67.7M	1e-3	48	50x100	01:49now	.422	.132	.385	.351	.111	.345	.351	.110	.335	.789	.238	.572	
175	3	1024	68.1M	1e-3	48	100x100	01:35	.277	.086	.270	.241	.075	.236	.246	.076	.242	.650	.193	.476	kernel size back to 5, this is the best model

I claim a 0.189 test-clean WER, please check output.json, the output of test.py.
More precisely, I have:

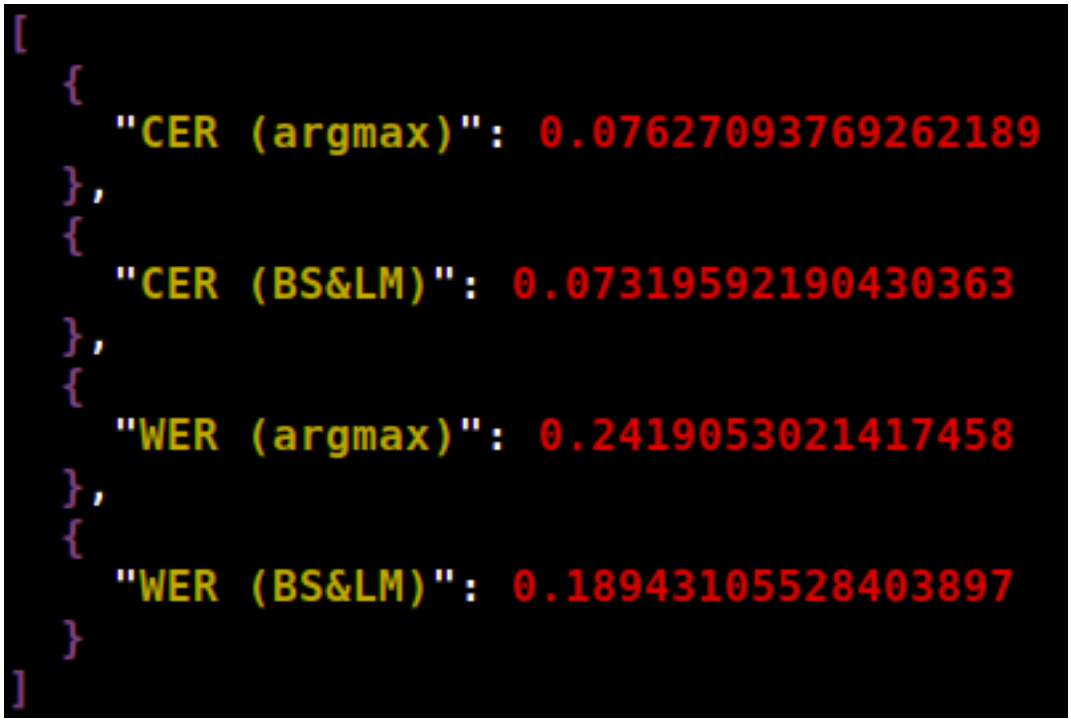

Please check the graphs and all the statistics of my runs in wandb, the relevant ones are listen at the first picture.
If you want to reproduce and run train.py and test.py yourself, please check the readme file.